

CSC311 Final Report

Fall 2020

Chunjing Zhang, Peiqing Yu


University of Toronto
2020.12.15

Part A

Question 1

(a) For $k = \{1, 6, 11, 16, 21, 26\}$, the corresponding validation accuracy imputed by user are:

```
Knn imputed by user have accuracy  
[0.6244707874682472, 0.6780976573525261, 0.6895286480383855, 0.6755574372001129, 0.6692068868190799, 0.6522720858029918]
```



Please check the relevant code at [knn.py](#)


(b) We report the test accuracy with k^* having the highest performance on validation data:

```
For knn imputed by users,  $k^* = 11$  has the highest performance on validation data.  
The final test accuracy is 0.6841659610499576.
```

(c) For $k = \{1, 6, 11, 16, 21, 26\}$, the corresponding validation accuracy imputed by item are:

```
Knn imputed by question have accuracy  
[0.607112616426757, 0.6542478125882021, 0.6826136042901496, 0.6860005644933672, 0.6922099915325995, 0.69037538808919]
```

```
For knn imputed by questions,  $k^* = 21$  has the highest performance on validation data.  
The final test accuracy is 0.6816257408975445.
```



- (d) Comparing the test performance between user and item based collaborative filtering, we can see that user-based imputation has higher test accuracy of 68.42% and item-based imputation has slightly lower test accuracy of 68.16%. Thus user-based collaborative filtering method performs better.
- (e) Limitation of KNN includes: slow computation with lots of memory occupied by training data
Relatively low prediction

Question 2

- (a) We can first derive the probability that question j is wrongly answered by student i:

$$p(c_{ij} = 0 | \theta_i, \beta_j) = 1 - \frac{\exp(\theta_i - \beta_j)}{1 + \exp(\theta_i - \beta_j)} = \frac{1}{1 + e^{\theta_i - \beta_j}}$$

$$p(c_{ij} = 1 | \theta_i, \beta_j) = \frac{\exp(\theta_i - \beta_j)}{1 + \exp(\theta_i - \beta_j)} = \frac{e^{\theta_i}}{e^{\theta_i} + e^{\beta_j}} = \frac{1}{1 + e^{\beta_j - \theta_i}}$$

Thus the likelihood of the whole training data is:

$$\prod_{k=1}^{56688} \frac{1}{1 + e^{\beta_{j(k)} - \theta_{i(k)}}}^{c_{ij(k)}} \frac{1}{1 + e^{\theta_{i(k)} - \beta_{j(k)}}}^{1 - c_{ij(k)}}$$

And we can derive the log-likelihood $\log p(C|\theta, \beta)$ for all students and questions:

$$\sum_{k=1}^{56688} c_{ij(k)} \log(1 + e^{\beta_{j(k)} - \theta_{i(k)}}) + (1 - c_{ij(k)}) \log(1 + e^{\theta_{i(k)} - \beta_{j(k)}})$$

The derivative of the log-likelihood with respect to θ_i and β_j are:

$$\frac{dl}{d\theta_i} = \sum_{k=1}^{56688} \frac{e^{\theta_{i(k)}}}{e^{\beta_{j(k)} + e^{\theta_{i(k)}}}} - c_{ij(k)}$$

$$\frac{dl}{d\beta_j} = \sum_{k=1}^{56688} c_{ij(k)} - \frac{e^{\theta_{i(k)}}}{e^{\beta_{j(k)} + e^{\theta_{i(k)}}}}$$

- (b) Please see the relevant code at item_response.py


```
For learning rate = 0.1, iteration = 50, the validation accuracy is
[0.48701665255433246, 0.6905165114309907, 0.5819926615862263, 0.6655376799322608, 0.6577758961332204, 0.6364662715213096, 0.6549534292972058, 0.6566469093988145, 0.6299745977984759, 0.6747106971493085, 0.6481795088907706, 0.647897262207169, 0.6501552356759808, 0.6676545300592718, 0.6315269545582839, 0.663279706463449, 0.6543889359300028, 0.647897262207169, 0.6423934518769404, 0.6699125035280835, 0.6350550381033023, 0.6652554332486593, 0.6486028789161727, 0.6521309624611911, 0.6501552356759808, 0.6600338696020321, 0.6414055884843353, 0.65961049957663, 0.6572114207660175, 0.6453570420547559, 0.6466514253457521, 0.6718882303132938, 0.6322325712672876, 0.6622918430708439, 0.6541066892464014, 0.6517075924357889, 0.6477561388653683, 0.6689246401354784, 0.6343494213942986, 0.6622918430708439, 0.658058142816822, 0.642675698560542, 0.6512842224103866, 0.663279706463449, 0.6385831216483207, 0.6615862263618403, 0.6474738921817669, 0.6543889359300028, 0.6529777025119955, 0.6622918430708439]
```

```
For learning rate = 0.01, iteration = 50, the validation accuracy is
[0.5033869602032176, 0.635196161445103, 0.6610217329946373, 0.6742873271239063, 0.6817668642393452, 0.
0.6875529212531752, 0.6902342647473892, 0.694044594976009, 0.6979960485464296, 0.7003951453570421, 0.
0.7059522156364663, 0.7056167090036692, 0.7061812023708721, 0.7071690657634773, 0.7084391758396839, 0.
0.7074513124470787, 0.7080158058142817, 0.707310189105278, 0.7064634490544736, 0.70575783234547, 0.
0.7053344623200677, 0.7050522156364663, 0.7044877222692634, 0.7040643522438611, 0.7050522156364663, 0.
0.7044877222692634, 0.7043465989274627, 0.7042054755856618, 0.7040643522438611, 0.7037821055602597, 0.
0.7039232289020604, 0.7042954755856618, 0.7037821055602597, 0.7047699689528648, 0.7058989556872707, 0.
0.7060400790290714, 0.7064634490544736, 0.7061812023708721, 0.7060400790290714, 0.7058989556872707, 0.
0.7060400790290714, 0.7061812023708721, 0.7060400790290714, 0.7056167090036692, 0.7054755856618685, 0.
0.7049110922946655, 0.7046288456110641, 0.7047699689528648, 0.7049110922946655, 0.7050522156364663]
```

```
For learning rate = 0.01, iteration = 100, the validation accuracy is
[0.49562517640417725, 0.6450747953711544, 0.6707592435788879, 0.6819079875811459, 0.6902342647473892, 0.6960203217612193, 0.7001128986734406, 0.7020886254586508, 0.7036409822184589, 0.7047699689528648, 0.707874682472481, 0.7075924357888794, 0.7071690657634773, 0.7077335591306803, 0.7090036692086868, 0.7098504092576913, 0.710132265594192927, 0.7104149026248942, 0.71069714933084956, 0.7098504092576913, 0.7085802991814846, 0.707874682472481, 0.7074513124470787, 0.707310189105278, 0.7068868190798758, 0.7074513124470787, 0.707874682472481, 0.7081569291560824, 0.707874682472481, 0.7077335591306803, 0.7070279424216765, 0.7067456957380751, 0.7070279424216765, 0.7064634490544736, 0.7058989556872707, 0.7061812023708721, 0.7060400790290714, 0.7061812023708721, 0.7064634490544736, 0.7066045723962744, 0.7067456957380751, 0.7063223257126728, 0.7061812023708721, 0.7061812023708721, 0.706400790290714, 0.7058989556872707, 0.7063223257126728, 0.7061812023708721, 0.7063223257126728, 0.7067456957380751, 0.7070279424216765, 0.7066045723962744, 0.7064634490544736, 0.7063223257126728, 0.7063223257126728, 0.7057583234547, 0.7061812023708721, 0.7058989556872707, 0.7058989556872707, 0.7057583234547, 0.7058989556872707, 0.7063223257126728, 0.7066045723962744, 0.7064634490544736, 0.7064634490544736, 0.7063223257126728, 0.7063223257126728, 0.7066045723962744, 0.7067456957380751, 0.7063223257126728, 0.7064634490544736, 0.7063223257126728, 0.7057583234547, 0.7057583234547, 0.7054755856618685, 0.7054755856618685, 0.7056167090036692, 0.7058989556872707, 0.7056167090036692, 0.7057583234547, 0.7058989556872707, 0.7058989556872707, 0.7056167090036692, 0.7056167090036692, 0.7056167090036692, 0.7056167090036692, 0.7053344623200677, 0.7056167090036692, 0.7054755856618685, 0.7056167090036692, 0.7056167090036692, 0.7056167090036692, 0.7054755856618685, 0.7058989556872707]
```

```
For learning rate = 0.015, iteration = 50, the validation accuracy is
[0.49689528648038384, 0.6587637595258256, 0.6782387806943269, 0.6900931414055885, 0.698278295230031, 0
.703076488851256, 0.7046288456110641, 0.7042054755856618, 0.7044877222692634, 0.7068868190798758, 0
.7053344623200677, 0.7063223257126728, 0.7061812023708721, 0.7060400790290714, 0.7064634490544736, 0
.70575783234547, 0.705616790036692, 0.7058989556872707, 0.7061812023708721, 0.7061812023708721, 0
.7058989556872707, 0.7061812023708721, 0.7058989556872707, 0.70575783234547, 0.7061812023708721, 0
.7061812023708721, 0.7063223257126728, 0.7060400790290714, 0.70575783234547, 0.70575783234547, 0.705616790036692,
0.7066045723962744, 0.7066045723962744, 0.7067456957380751, 0.7068868190798758, 0.7067456957380751, 0
.7064634490544736, 0.7064634490544736, 0.7063223257126728, 0.7064634490544736, 0.7058989556872707, 0
.7060400790290714, 0.7058989556872707, 0.7064634490544736, 0.7060400790290714, 0.7060400790290714, 0
.70575783234547, 0.70575783234547, 0.7054755856618685, 0.7053344623200677]
```


Here is the training curve that shows the training and validation log-likelihoods as a function of iteration with learning rate = 0.15, iteration = 50



(c) The final test and validation accuracy are:

```
With chosen learning rate = 0.015 and iteration = 50,
The final validation accuracy is 0.7053344623200677
The final test accuracy is 0.7064634490544736
```

(d) The probability of the correct response of question # 0, 1, 2, 3, 4 are:



We can see that the shape of the curves is sigmoid. This shape indicates that as the student ability increases, the probability of correctly answered the specific questions increases. This shape is consistent with all the randomly chosen question.

3

Option 2: Neural Networks.

While running nn, we got the UserWarning: nn.functional.sigmoid is deprecated. Use torch.sigmoid instead. However these two functions are identical, so we keep nn.functional.sigmoid, which is provided in the starter code.

(a)

ALS (Alternating Least Squares) is used in unsupervised learning, while neural networks is used for supervised learning.

ALS is a method used in a matrix multiplication to get to the final answer. While neural networks can be used more widely in other situations (universality).

ALS is updating one parameter at a time while fixing the other ones and alternating between the parameters, while neural networks update all the parameters at the same time using a matrix form.

(b)

Codes are in py file.

(c)

Test learning rate $\{0.01, 0.05, 0.1, 0.5\}$ and num_epoch under 30 (mainly around 10) for each k . The best validation accuracy with corresponding hyperparameters are below:

k	lr	num_epoch	accuracy
10	0.1	10	0.6862828
50	0.05	10	0.6868473
100	0.05	7	0.6803556
200	0.05	9	0.6765453
500	0.1	9	0.6732995

According to results above, $k^* = 50$.

$k = 10$:

Epoch: 0	Training Cost: 13509.171195	Valid Acc: 0.6193903471634208
Epoch: 1	Training Cost: 12262.740651	Valid Acc: 0.6395709850409258
Epoch: 2	Training Cost: 11545.053267	Valid Acc: 0.658058142816822
Epoch: 3	Training Cost: 10942.724810	Valid Acc: 0.6670900366920689
Epoch: 4	Training Cost: 10450.880671	Valid Acc: 0.6724527236804968
Epoch: 5	Training Cost: 10044.467480	Valid Acc: 0.68077900884674
Epoch: 6	Training Cost: 9703.641267	Valid Acc: 0.6817668642393452
Epoch: 7	Training Cost: 9410.230672	Valid Acc: 0.6814846175557437
Epoch: 8	Training Cost: 9142.778405	Valid Acc: 0.6830837143663562
Epoch: 9	Training Cost: 8909.321104	Valid Acc: 0.6862828111769687

$k = 50$:

Epoch: 0	Training Cost: 13340.221335	Valid Acc: 0.6271521309624611
Epoch: 1	Training Cost: 12292.556435	Valid Acc: 0.642111205193339
Epoch: 2	Training Cost: 11588.957260	Valid Acc: 0.6510019757267852
Epoch: 3	Training Cost: 10923.860900	Valid Acc: 0.665961049957663
Epoch: 4	Training Cost: 10311.696580	Valid Acc: 0.6752751905165114
Epoch: 5	Training Cost: 9753.755806	Valid Acc: 0.6804967541631386
Epoch: 6	Training Cost: 9235.304152	Valid Acc: 0.6867061812023709
Epoch: 7	Training Cost: 8743.937502	Valid Acc: 0.6874117979113745
Epoch: 8	Training Cost: 8274.097998	Valid Acc: 0.6869884278859724
Epoch: 9	Training Cost: 7824.574489	Valid Acc: 0.6868473045441716

$k = 100$:

Epoch: 0	Training Cost: 13454.517941	Valid Acc: 0.6255997742026531
Epoch: 1	Training Cost: 12410.300181	Valid Acc: 0.6416878351679368
Epoch: 2	Training Cost: 11614.630996	Valid Acc: 0.6589048828676263
Epoch: 3	Training Cost: 10856.562655	Valid Acc: 0.6661021732994638
Epoch: 4	Training Cost: 10151.541637	Valid Acc: 0.6756985605419137
Epoch: 5	Training Cost: 9480.503721	Valid Acc: 0.6776742873271239
Epoch: 6	Training Cost: 8818.394983	Valid Acc: 0.6803556308213379

$k = 200$:


Epoch: 0	Training Cost: 13807.696970	Valid Acc: 0.6169912503528083
Epoch: 1	Training Cost: 12684.508000	Valid Acc: 0.6354784081287045
Epoch: 2	Training Cost: 11820.206228	Valid Acc: 0.6487440022579735
Epoch: 3	Training Cost: 11065.518775	Valid Acc: 0.6598927462602314
Epoch: 4	Training Cost: 10352.613084	Valid Acc: 0.6642675698560542
Epoch: 5	Training Cost: 9632.534934	Valid Acc: 0.6732994637313011
Epoch: 6	Training Cost: 8881.009156	Valid Acc: 0.6714648602878917
Epoch: 7	Training Cost: 8100.275627	Valid Acc: 0.6735817104149027
Epoch: 8	Training Cost: 7312.334899	Valid Acc: 0.676545300592718

$k = 500$:


Epoch: 0	Training Cost: 15355.295413	Valid Acc: 0.6049957662997459
Epoch: 1	Training Cost: 13455.437402	Valid Acc: 0.6240474174428451
Epoch: 2	Training Cost: 12329.551403	Valid Acc: 0.6346316680779001
Epoch: 3	Training Cost: 11443.095316	Valid Acc: 0.6480383855489698
Epoch: 4	Training Cost: 10657.886461	Valid Acc: 0.6573525261078102
Epoch: 5	Training Cost: 9837.048679	Valid Acc: 0.6658199266158623
Epoch: 6	Training Cost: 8894.015684	Valid Acc: 0.672311600338696
Epoch: 7	Training Cost: 7817.088929	Valid Acc: 0.6732994637313011
Epoch: 8	Training Cost: 6677.067158	Valid Acc: 0.6732994637313011

(d)

Need to plot training loss, validation loss, and validation accuracy.
Validation accuracy and epoch:



Train & validation loss and epoch:



The validation accuracy increases as epoch increases. While the loss for both train and validation decrease as epoch increases. However, the MSE for validation decrease more and more slowly, nearly becomes flat at the end.

(e)

For $k^* = 50$, $lr = 0.05$, $num_epoch = 10$, got this for each λ in list $\{0.001, 0.01, 0.1, 1\}$:

λ	validation acc	test acc
0.001	0.6879763	0.6838837
0.01	0.6860006	0.6771098
0.1	0.6907988	0.6827547
1	0.6241885	0.6226362

Here are results for each λ :

$\lambda = 0.001$:

Epoch: 0	Training Cost: 13329.817492	Valid Acc: 0.6308213378492803
Epoch: 1	Training Cost: 12269.545421	Valid Acc: 0.6406999717753317
Epoch: 2	Training Cost: 11546.206075	Valid Acc: 0.6584815128422241
Epoch: 3	Training Cost: 10883.661262	Valid Acc: 0.6679367767428732
Epoch: 4	Training Cost: 10292.029622	Valid Acc: 0.6752751905165114
Epoch: 5	Training Cost: 9758.018521	Valid Acc: 0.679226644086932
Epoch: 6	Training Cost: 9258.252411	Valid Acc: 0.682895850973751
Epoch: 7	Training Cost: 8780.376137	Valid Acc: 0.6855771944679651
Epoch: 8	Training Cost: 8319.096074	Valid Acc: 0.687694044594976
Epoch: 9	Training Cost: 7875.928446	Valid Acc: 0.6879762912785775
0.6838837143663562		

$\lambda = 0.01$:

Epoch: 0	Training Cost: 13432.491016	Valid Acc: 0.6275755009878634
Epoch: 1	Training Cost: 12384.960312	Valid Acc: 0.6392887383573242
Epoch: 2	Training Cost: 11696.357103	Valid Acc: 0.6556590460062094
Epoch: 3	Training Cost: 11050.797025	Valid Acc: 0.6662432966412645
Epoch: 4	Training Cost: 10461.818962	Valid Acc: 0.6725938470222975
Epoch: 5	Training Cost: 9934.681867	Valid Acc: 0.6776742873271239
Epoch: 6	Training Cost: 9448.708539	Valid Acc: 0.6806378775049393
Epoch: 7	Training Cost: 8986.494424	Valid Acc: 0.6830369743155518
Epoch: 8	Training Cost: 8542.530875	Valid Acc: 0.6861416878351679
Epoch: 9	Training Cost: 8116.903748	Valid Acc: 0.6860005644933672
0.677109793959921		

$\lambda = 0.1$:

Epoch: 0	Training Cost: 14133.327981	Valid Acc: 0.6257408975444538
Epoch: 1	Training Cost: 13232.630282	Valid Acc: 0.6394298616991251
Epoch: 2	Training Cost: 12706.057663	Valid Acc: 0.6521309624611911
Epoch: 3	Training Cost: 12237.017430	Valid Acc: 0.661727349703641
Epoch: 4	Training Cost: 11830.097083	Valid Acc: 0.6696302568444821
Epoch: 5	Training Cost: 11480.480174	Valid Acc: 0.6742873271239063
Epoch: 6	Training Cost: 11171.502251	Valid Acc: 0.6817668642393452
Epoch: 7	Training Cost: 10886.335584	Valid Acc: 0.6840715777589613
Epoch: 8	Training Cost: 10614.760466	Valid Acc: 0.6869884278859724
Epoch: 9	Training Cost: 10352.705603	Valid Acc: 0.6907987581145921

0.6827547276319503

$\lambda = 1$:

Epoch: 0	Training Cost: 18829.117218	Valid Acc: 0.6233418007338414
Epoch: 1	Training Cost: 15533.271581	Valid Acc: 0.6182613604290149
Epoch: 2	Training Cost: 14301.143335	Valid Acc: 0.6147332768039966
Epoch: 3	Training Cost: 14020.351263	Valid Acc: 0.6168501270110076
Epoch: 4	Training Cost: 13803.935659	Valid Acc: 0.619813717188823
Epoch: 5	Training Cost: 13602.622974	Valid Acc: 0.6213660739486311
Epoch: 6	Training Cost: 13421.036527	Valid Acc: 0.6212249506068304
Epoch: 7	Training Cost: 13263.540402	Valid Acc: 0.6232006773920407
Epoch: 8	Training Cost: 13131.815360	Valid Acc: 0.6230595540502399
Epoch: 9	Training Cost: 13023.678681	Valid Acc: 0.6241885407846458

0.6226361840248377

It seems like the model performs slightly better for $\lambda = 0.1$, which makes validation accuracy reach 0.69.

For this question, we used the neural networks as base model. First, we used `np.random.randint()` to randomly select n students with replacement for three times, where n is the total number of students. Then trained three base models corresponding to the three groups of selected data (selected rows from the data matrix where each row represents a selected students with all questionss). Then get prediction by the three models' average for evaluation.

```
Epoch: 0    Training Cost: 13571.452437    Valid Acc: 0.6079593564775614
Epoch: 1    Training Cost: 12145.219665    Valid Acc: 0.6038667795653401
Epoch: 2    Training Cost: 11077.611763    Valid Acc: 0.5970928591589049
Epoch: 3    Training Cost: 10169.012415    Valid Acc: 0.5906011854360711
Epoch: 4    Training Cost: 9452.807407     Valid Acc: 0.581145921535422
Epoch: 5    Training Cost: 8879.751910     Valid Acc: 0.581145921535422
Epoch: 6    Training Cost: 8407.603848     Valid Acc: 0.5817104149026249
Epoch: 7    Training Cost: 8010.175925     Valid Acc: 0.5798758114592154
Epoch: 8    Training Cost: 7670.466305     Valid Acc: 0.5786057013830087
Epoch: 9    Training Cost: 7376.891085     Valid Acc: 0.578464578041208
Epoch: 0    Training Cost: 13467.668539    Valid Acc: 0.6011854360711262
Epoch: 1    Training Cost: 12387.411854    Valid Acc: 0.6007620660457239
Epoch: 2    Training Cost: 11526.810686    Valid Acc: 0.596528365791702
Epoch: 3    Training Cost: 10690.442652    Valid Acc: 0.5930002822466836
Epoch: 4    Training Cost: 9994.157496     Valid Acc: 0.586931978549252
Epoch: 5    Training Cost: 9415.572773     Valid Acc: 0.5824160316116286
Epoch: 6    Training Cost: 8923.523271     Valid Acc: 0.5776178379904036
Epoch: 7    Training Cost: 8501.869490     Valid Acc: 0.5771944679650014
Epoch: 8    Training Cost: 8138.308293     Valid Acc: 0.5767710979395992
Epoch: 9    Training Cost: 7822.550102     Valid Acc: 0.574795371154389
Epoch: 0    Training Cost: 14249.315764    Valid Acc: 0.6079593564775614
Epoch: 1    Training Cost: 12882.573984    Valid Acc: 0.6104995766299746
Epoch: 2    Training Cost: 11897.499172    Valid Acc: 0.6064069997177534
Epoch: 3    Training Cost: 11055.499102    Valid Acc: 0.6034434095399379
Epoch: 4    Training Cost: 10356.602646    Valid Acc: 0.6013265594129269
Epoch: 5    Training Cost: 9763.409961     Valid Acc: 0.5992097092859159
Epoch: 6    Training Cost: 9253.768347     Valid Acc: 0.599633079311318
Epoch: 7    Training Cost: 8810.833066     Valid Acc: 0.5997742026531189
Epoch: 8    Training Cost: 8423.033213     Valid Acc: 0.5977984758679086
Epoch: 9    Training Cost: 8082.173427     Valid Acc: 0.5953993790572961
0.5935647756138865
0.5901778154106689
```

The final test accuracy is 0.5901778154106689 which is much lower than the single neural networks model. The final validation accuracy is 0.5935647756138865, also much lower than before.

Therefore, actually we didn't obtain better performance using the ensemble, and the running time for running ensemble once is quite long.

We think one main reason why accuracy decline could because when sampling with replacement, some data is not sampled into all of the three samples. Also, some model with wrong prediction may have higher weight for its wrong prediction which lead to a wrong results finally.

Part B

1 Formal Description


We have tried multiple ways to modify the implemented algorithm in part A to increase the test accuracy on students' answers to the diagnostic questions. Above all, we tried to consider the improvement on the neural networks algorithm. Since tuning hyperparameter is already done in Part A, we tried something else. For example, first we tried to add one hidden layer. Then we tried to use other activation functions like ReLU, Softmax, and LogSigmoid. Also, we tried different regularizer (ie. L_1 regularization). Finally we also thought about normalize the input data; however, since the input data is not in different scale we didn't modify the normalization.

After examining all the possible modifications on the neural network algorithm, we took a deeper look at the the student metadata, `student_meta.csv`. This dataset is composed of the user id, gender, date of birth and premium pupil. As the training accuracy of the baseline algorithms in part A is not too good, we hypothesized that there might exists a mixture of data where each group of data has its own density for the provided training, validation and test dataset. If we train the model with the combined data, the model may be weak to capture the overall existed pattern, and this may cause underfitting. The gender and age of users may both be the factor that affect the probability of answering questions correctly. This makes sense as younger students normally have lower accuracy, and as children get older they tend to have higher accuracy. Therefore, we separate the training, validation and test data into individual dictionaries based on gender and age, then train the item-response probability algorithm on each of them.

The separation of dataset based on age is at `age.py` and the separation of dataset based on gender is at `gender.py`. We expect the separate training may increase the train accuracy and reduce underfit, and optimally increase the test accuracy.

2 Figure or Diagram

Neural Network Modification



Add a hidden layer, also update the regularizer:

```
# Define linear functions.  
self.g = nn.Linear(num_question, k)  
self.s = nn.Linear(k, k)  
self.h = nn.Linear(k, num_question)
```

```
g_w_norm = torch.norm(self.g.weight, 2)
s_w_norm = torch.norm(self.s.weight, 2)
h_w_norm = torch.norm(self.h.weight, 2)
```

Tried different activation functions with the added new layer:

```
activate_g = self.g(inputs)
output1 = F.sigmoid(activate_g)
# output = torch.sigmoid(activate_g)
# m = nn.ReLU()
# m = nn.LogSigmoid()
# m = nn.LogSoftmax()
# output = m(activate_g)
activate_s = self.s(output1)
output2 = F.sigmoid(activate_s)
activate_f = self.h(output2)
out = F.sigmoid(activate_f)
# out = torch.sigmoid(activate_f)
# out = m(activate_f)
```

Group separation using Student Metadata

This is the code snippets implemented the separate training of each group data based on gender information provided in the student metadata.

```
# separate the train, validation, and test data by gender
boy, girl = separate_gender()
b_test, g_test = get_dic(boy, girl, test_data)
b_train, g_train = get_dic(boy, girl, train_data)
b_val, g_val = get_dic(boy, girl, val_data)

# train each male and female datasets separately
boy = irt(b_train, b_val, 0.015, 50)
girl = irt(g_train, g_val, 0.015, 50)
print(f"For learning rate = 0.015, iteration = 50, the train accuracy for boys is\n{boy[2]}")
print(f"For learning rate = 0.015, iteration = 50, the train accuracy for girls is\n{girl[2]}")
```


3 Comparison or Demonstration

Neural Network Modification

Here are results for using Sigmoid with an additional hidden layer:

Epoch: 0	Training Cost: 14173.529993	Valid Acc: 0.6277166243296641
Epoch: 1	Training Cost: 13453.212581	Valid Acc: 0.6305390911656789
Epoch: 2	Training Cost: 13106.100574	Valid Acc: 0.6349139147615016
Epoch: 3	Training Cost: 12803.304464	Valid Acc: 0.6412644651425345
Epoch: 4	Training Cost: 12560.058697	Valid Acc: 0.648461755574372
Epoch: 5	Training Cost: 12356.876511	Valid Acc: 0.6549534292972058
Epoch: 6	Training Cost: 12266.673736	Valid Acc: 0.6668077900084673
Epoch: 7	Training Cost: 12193.919032	Valid Acc: 0.668077900084674
Epoch: 8	Training Cost: 12119.076354	Valid Acc: 0.668077900084674
Epoch: 9	Training Cost: 12077.401896	Valid Acc: 0.6737228337567034
Epoch: 10	Training Cost: 12133.146861	Valid Acc: 0.6721704769968952
Epoch: 11	Training Cost: 12035.328566	Valid Acc: 0.6747106971493085
Epoch: 12	Training Cost: 11957.835894	Valid Acc: 0.6776742873271239
Epoch: 13	Training Cost: 11905.195812	Valid Acc: 0.677109793959921
Epoch: 14	Training Cost: 11822.732113	Valid Acc: 0.6806378775049393
Epoch: 15	Training Cost: 11807.365512	Valid Acc: 0.6779565340107254
Epoch: 16	Training Cost: 11770.325193	Valid Acc: 0.6827547276319503
Epoch: 17	Training Cost: 11752.178491	Valid Acc: 0.6816257408975445
Epoch: 18	Training Cost: 11741.590530	Valid Acc: 0.682895850973751
Epoch: 19	Training Cost: 11678.233868	Valid Acc: 0.6844482077335591
Epoch: 20	Training Cost: 11618.268810	Valid Acc: 0.6833192209991532
Epoch: 21	Training Cost: 11622.833282	Valid Acc: 0.6827547276319503
Epoch: 22	Training Cost: 11655.201063	Valid Acc: 0.6864239345187694
Epoch: 23	Training Cost: 11599.434348	Valid Acc: 0.6857183178097658
Epoch: 24	Training Cost: 11532.738863	Valid Acc: 0.6850127011007621
		0.6782387806943269
		0.6850127011007621

Where final validation accuracy is 0.6850127011007621 and test accuracy is 0.6782387806943269.




However, there's no obvious changing for using this.

Then by using ReLU as activation function, we got this:

```
Epoch: 0   Training Cost: 43074.088623   Valid Acc: 0.5193338978267006
Epoch: 1   Training Cost: 35417.084000   Valid Acc: 0.5397967823878069
Epoch: 2   Training Cost: 33547.767464   Valid Acc: 0.5438893593000282
Epoch: 3   Training Cost: 33337.407709   Valid Acc: 0.5467118261360429
Epoch: 4   Training Cost: 32972.840151   Valid Acc: 0.5496754163138583
Epoch: 5   Training Cost: 32852.308329   Valid Acc: 0.5491109229466554
Epoch: 6   Training Cost: 32645.523170   Valid Acc: 0.5517922664408693
Epoch: 7   Training Cost: 32485.007030   Valid Acc: 0.5534857465424782
Epoch: 8   Training Cost: 32290.693451   Valid Acc: 0.5532034998588766
Epoch: 9   Training Cost: 32172.030867   Valid Acc: 0.554614733276884
0.563646627152131
```


Validation accuracy is 0.554614733276884 and test accuracy is 0.563646627152131.



This also doesn't help, since ReLU is not a better activation function compared to Sigmoid. We then tried L_1 regularization:

```
Epoch: 0    Training Cost: 27082.602104    Valid Acc: 0.6123341800733841
Epoch: 1    Training Cost: 20425.982297    Valid Acc: 0.6233418007338414
Epoch: 2    Training Cost: 20050.000346    Valid Acc: 0.6253175275190517
Epoch: 3    Training Cost: 19795.116239    Valid Acc: 0.6270110076206604
Epoch: 4    Training Cost: 19618.048839    Valid Acc: 0.6270110076206604
Epoch: 5    Training Cost: 19489.839846    Valid Acc: 0.6275755009878634
Epoch: 6    Training Cost: 19395.148801    Valid Acc: 0.626728760937059
Epoch: 7    Training Cost: 19324.223396    Valid Acc: 0.6264465142534575
Epoch: 8    Training Cost: 19267.709228    Valid Acc: 0.6258820208862546
Epoch: 9    Training Cost: 19223.142982    Valid Acc: 0.6254586508608524
0.6223539373412362
```

Validation accuracy is 0.6254586508608524 and test accuracy is 0.6223539373412362.




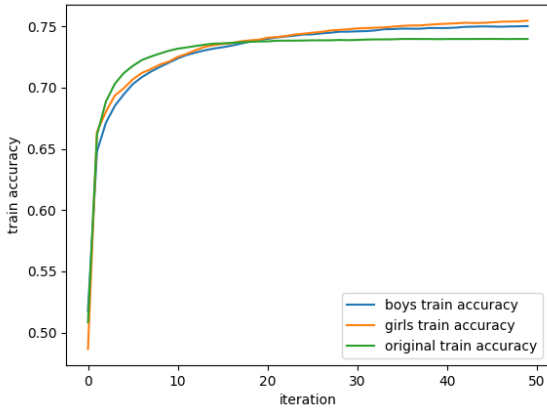
Similarly, this doesn't help to improve our results.

Group separation using Student Metadata

We can see that there is significant increase of training accuracy for all the gender groups. We can observe that the highest train accuracy for boy group has reached 75%. However, the validation accuracy have shown no improvement. The test accuracy of one group increases 2% and the other group stays the same.

```
For learning rate = 0.015, iteration = 50, the train accuracy for boys is
[0.5176057984568623, 0.6478372691138649, 0.6714051905541267, 0.6852466682253916, 0.6946457797521627, 0
.7029693710544774, 0.7085807809212065, 0.7130699088145896, 0.7169043722235212, 0.7203647416413373, 0
.723871872808043, 0.7267243394902969, 0.7286883329436521, 0.7305120411503391, 0.7319616553659107, 0
.7330371755903671, 0.7343465045592705, 0.7359831657703998, 0.7375263034837503, 0.7389291559504325, 0
.7398643909282208, 0.7409866729015665, 0.7417816226326864, 0.7424830488660276, 0.7434182838438158, 0
.743558569090484, 0.7443535188216039, 0.7449146598082769, 0.7457096095393968, 0.7457563712882862, 0
.7460369417816226, 0.7462239887771802, 0.7467851297638531, 0.7477203647416414, 0.747907411737199, 0
.7482815057283142, 0.7481879822305354, 0.7481879822305354, 0.7487491232172083, 0.7487491232172083, 0
.7487491232172083, 0.7490764554594341, 0.7496843581949965, 0.7498714051905542, 0.7500584521861118, 0
.749964928688333, 0.7498246434416648, 0.7500116904372224, 0.7501519756838906, 0.7502922609305588]
```

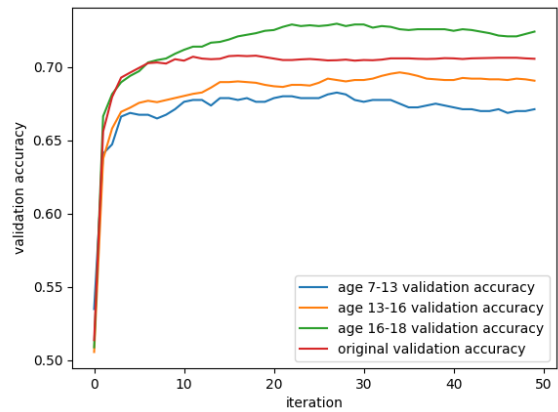
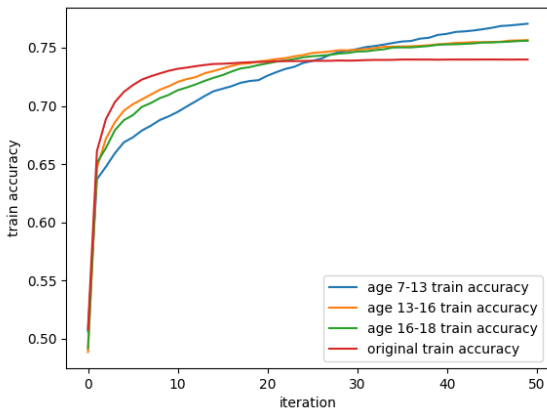
```
With chosen learning rate = 0.015 and iteration = 50,
The final test accuracy for age 7-13 is 0.6238532110091743
The final test accuracy for age 13-16 is 0.7026178010471205
The final test accuracy for age 16-18 is 0.7263479145473042
```

The train accuracy of all the age group separation increased to a higher extent, and we can observe that the highest train accuracy for age 7-13 has reached 76.8%. However, the validation accuracy and test accuracy have shown an overall no improvement.

```
For learning rate = 0.015, iteration = 50, the train accuracy for age 7-13 is
[0.4785867237687366, 0.637503823799327, 0.6494340776996024, 0.6563169164882227, 0.6648822269807281, 0.671612113796268, 0.6751300091771184, 0.6789538085041297, 0.6838482716427042, 0.6882838788620373, 0.6916488222698073, 0.6951667176506577, 0.6991434689507494, 0.7018966044661976, 0.7055674518201285, 0.7106148669317834, 0.7129091465279902, 0.7184154175588865, 0.7213215050474151, 0.7245334964821046, 0.7274395839706332, 0.7294279596206791, 0.7315692872438054, 0.7330988069746099, 0.7358519424900581, 0.737840318140104, 0.7396757418170694, 0.7416641174671154, 0.7438054450902417, 0.7454879167941266, 0.7470174365249311, 0.7486999082288162, 0.7503823799327012, 0.752370755582747, 0.7534414193943102, 0.7543591312327929, 0.755429795044356, 0.7561945549097583, 0.7565004588559192, 0.7578770266136433, 0.7601713062098501, 0.760477210156011, 0.7615478739675742, 0.7617008259406546, 0.7641480575099419, 0.7649128173753441, 0.7661364331599878, 0.767207096971551, 0.7670541449984705, 0.7678189048638727]
```

```
With chosen learning rate = 0.015 and iteration = 50,
The final test accuracy for age 7-13 is 0.6238532110091743
The final test accuracy for age 13-16 is 0.7026178010471205
The final test accuracy for age 16-18 is 0.7263479145473042
```



4 Limitations

- Adding hidden layer requires tuning for hyperparameters again. Since the original k may lead to overfitting, while the original learning rate and number of iterations may lead to relatively longer running time. However, adding 1 layer seems to help less, doesn't contribute to both accuracy and loss in our experiment.
- One possible extension could be, trying to add more hidden layers, then tune hyperparameters to avoid overfitting. We think this may help to improve accuracy. Also, maybe mixing all methods (ie. k , number of hidden layers, activation function) we tried to use could also bring improvements. Since there could be some better combinations if try more.
- The data processing for converting given train dataset into sparse matrix is time consuming and non-efficient.
- Even though we reduce the underfitting and increase train accuracy, the Group separation using student metadata is not modifying the algorithm, but focuses on data processing. The improved train accuracy would not apply to a new dataset.

Contribution

Group Member: Peiqing Yu

- Part A: `knn.py`
- Part A: `item_response.py`
- Part A: report question 1, question 2
- Part B: `irt_partB.py`
- Part B: `modified_nn.py`
- Part B: half of the report

Group Member: Chunjing Zhang

- Part A: `neural_network.py`
- Part A: `esemble.py`
- Part A: report question 3, question 4
- Part B: `age.py`
- Part B: `gender.py`
- Part B: half of the report